

$\mathcal{N} = 2$ theories

Mathematical background and implementation reference

One reference, organized by package

Exact Lie-algebra data, anomaly checks, superconformal indices,
Coulomb-branch indices and spectra, theory properties,
database persistence, and the planned Higgs-branch extension.

Database update: schema version 6

This update adds a database relationship diagram, a catalog of all ten tables and their keys, and the full/half-hyper normalization rules. It documents the $5 \rightarrow 6$ migration, conditional legacy-JSON cleanup, and the separate migration still needed for old realization hashes. The validation snapshot is refreshed.

Snapshot	Base reference: 8 September 2026. Database and validation update: 9 September 2026; 662ad31 plus working-tree changes.
Consolidates	<code>index/n2_theory_index_Mathematical_Background.pdf</code> and the previous <code>output/pdf/n2_implementation_reference_summary.pdf</code> .
Status	Database structure checked against the current source; other package chapters retain the base reference snapshot. Higgs/Hall–Littlewood material is explicitly marked as proposed, not implemented.
Editing scope	Documentation only. The legacy index-background PDF is retained; this consolidated reference supersedes it.

1 Project map and conventions

Package	Responsibility	Read here
<code>anomalies</code>	Cartan types, representations, full/half hypers, beta functions, local and conventional mod-two gauge anomalies.	§2
<code>index</code>	FORM expansion, LiE gauge projection and cache; Coulomb PE/PL and full-index projection; future Higgs/flavor refinement.	§3
<code>common</code>	Property orchestration, flavor blocks, central charges, exact-number and FORM helpers, JSON and MySQL persistence.	§4
<code>test</code>	Unit and backend regression tests, precision and serialization checks, optional live database integration.	§5

How data moves

Theory dictionary / JSON $\longrightarrow$ validated gauge and hyper data
 $\longrightarrow$ anomaly and conformality checks
 $\longrightarrow$ properties, full index, Coulomb index and spectrum
 $\longrightarrow$ exact JSON and optional MySQL storage

Standalone index functions are also callable directly. The property and database workflows impose the SCFT-candidate checks; a gauge-invariant counting calculation alone is not a proof that an input defines an SCFT.

Notation used throughout

$G = \prod_{a=1}^s G_a$	Compact, simply connected, semisimple gauge group; no abelian gauge factor or global quotient is implemented.
$\mathcal{R}_h = \boxtimes_a R_{ha}$	One product-group hyper representation; n_h is its number of copies. Omitted factors are singlets.
$\epsilon_h = 1, \frac{1}{2}$	Weight of a full or half hyper, respectively. Beta functions use $2\epsilon_h T$, not $\epsilon_h T$.
$E, R, r; j_1, j_2$	Scaling dimension, $SU(2)_R$ weight, $U(1)_r$ charge, and Lorentz Cartan weights in the chosen right-index convention.
t, y, u	Project full-index fugacities; the common literature variable $v = u^2$.
x, τ, q	Coulomb grading, proposed HL grading, and an internal FORM integer-power variable. This q is not a conventional full-index fugacity.

What “spectrum” means here

The Coulomb spectrum is the multiset of dimensions of independent chiral-ring generators, with repetitions. Index coefficients count all graded ring elements, including products. Neither is the spectrum of massive BPS particles or the full operator spectrum of the SCFT.

2 Package anomalies

2.1 Root data and representation invariants

MODULE `anomalies/lie_algebra.py`

`SimpleLieAlgebra` collects Sage root, weight and character objects together with the rank, dimension, dual Coxeter number and adjoint Dynkin labels. The cached constructor is `get_lie_algebra`. The accepted finite types are A_r ($r \geq 1$), B_r, C_r ($r \geq 2$), D_r ($r \geq 4$), $E_{6,7,8}$, F_4 , and G_2 . Use A_1 for B_1/C_1 , A_3 for D_3 , and $A_1 \times A_1$ for D_2 .

For a dominant integral highest weight λ , Dynkin labels and dimensions are

$$\lambda = \sum_i a_i \omega_i, \quad a_i = \langle \lambda, \alpha_i^\vee \rangle \in \mathbb{Z}_{\geq 0}, \quad (2.1)$$

$$\dim V_\lambda = \prod_{\alpha > 0} \frac{(\lambda + \rho, \alpha)}{(\rho, \alpha)}, \quad \rho = \frac{1}{2} \sum_{\alpha > 0} \alpha. \quad (2.2)$$

The algebra dimension is $\dim \mathfrak{g} = \text{rank } \mathfrak{g} + |\Phi|$, and the adjoint highest weight is the highest root. `representation_dimension` uses the degree of the Sage Weyl character, implementing (2.2). Source: Sage root-system and Weyl-character documentation [1, 2].

To compensate for the ambient root-length convention used by Sage, the code uses

$$C_2(\lambda) = \frac{(\lambda, \lambda + 2\rho)}{\ell_{\text{long}}^2}, \quad T(R) = \frac{\dim R}{\dim \mathfrak{g}} C_2(R). \quad (2.3)$$

If long roots have squared length two, the first denominator is two. Thus $C_2(\text{adj}) = T(\text{adj}) = h^\vee$ and $T(\mathbf{N}) = 1/2$ for $SU(N)$. The routines are `quadratic_casimir` and `dynkin_index`. The physics normalization is fixed by $\text{Tr}_R(T^a T^b) = T(R) \delta^{ab}$.

Names, conjugation and reality

`named_representation_labels` supports common names such as fundamental, antifundamental, adjoint, singlet, classical spinors and selected two-index representations. Explicit Dynkin labels are the general interface. Numbering is Sage’s finite-Cartan numbering; do not copy a label tuple from another reference without checking its node convention.

$$\lambda^* = -w_0 \lambda, \quad \nu(R) = \int_G \chi_R(g^2) d\mu(g) \in \{+1, 0, -1\}. \quad (2.4)$$

Here $\nu = +1, 0, -1$ means real, complex, pseudoreal. `conjugate_dynkin_labels` uses the dual character; `representation_reality` uses its Frobenius–Schur indicator. For an irreducible external tensor product, $\nu(\boxtimes_a R_a) = \prod_a \nu(R_a)$. Source: Sage Weyl characters [2]; representation and index conventions can be cross-checked with LieART 2.0 [4].

2.2 Theory data, hypermultiplets and the beta function

MODULE `anomalies/check_n2_anomalies.py`

The general input parser constructs `GaugeFactorData(factor_id, algebra)` objects and simple/product hyper data. Counts are exact nonnegative integers. Malformed types, labels or representations are rejected before calculation. The theory dictionary supports either one simple factor or a list of named gauge factors with a representation for each charged factor.

In $\mathcal{N} = 1$ language, a full hyper in $\mathcal{R}$ contains chirals in $\mathcal{R} \oplus \overline{\mathcal{R}}$. A half hyper contains one chiral and is allowed only when the *entire* gauge representation is pseudoreal. Consequently $\mathbf{2} \boxtimes \mathbf{2}$ of $SU(2)^2$ is real and cannot be a single half hyper, whereas $\mathbf{2} \boxtimes \mathbf{2} \boxtimes \mathbf{2}$ is pseudoreal and can. This is validated rather than inferred from one factor alone. Source: Bhardwaj–Tachikawa, §2, Eqs. (2.1)–(2.4) [3].

For factor a , spectator dimensions multiply the effective number of charged degrees of freedom. In the normalization of (2.3),

$$b_{0,a} = 2h_a^\vee - \sum_h n_h (2\epsilon_h) T_a(R_{ha}) \prod_{b \neq a} \dim R_{hb}. \quad (2.5)$$

The implementation uses the literal coefficient two for full hypers and one for half hypers. $b_{0,a} = 0$ for every simple factor is its one-loop conformality condition; $b_{0,a} > 0$ denotes asymptotic freedom in this convention. Source: Bhardwaj–Tachikawa, Eqs. (2.5)–(2.8) [3].

Checks that expose normalization mistakes

For $SU(N)$ with N_f full fundamental hypers,

$$b_0 = 2N - N_f. \quad (2.6)$$

For one full bifundamental $(\mathbf{N}, \mathbf{M})$ of $SU(N) \times SU(M)$, its contributions to the two beta-function subtractions are M and N . For a half hyper in an $SU(2)$ fundamental, the subtraction is $1/2$. These follow directly from (2.5), not from a new convention for T .

Two different uses of “index”

The Dynkin index $T(R)$ is a quadratic trace invariant. It is unrelated to the superconformal index $\mathcal{I}$. LieART’s commonly tabulated quadratic index $\ell(R)$ in the old project comparison is $2T(R)$; convert the convention before comparing tables. Root/node numbering must be converted independently.

Scope of the candidate flag

The anomaly checker combines valid matter, gauge-anomaly cancellation and vanishing beta functions into a `lagrangian_scft_candidate` result. It does not reconstruct a Seiberg–Witten geometry, identify S-dual theories, classify line operators, or handle non-Lagrangian matter sectors.

2.3 Local and conventional global gauge anomalies

For a left-handed Weyl fermion in R , the local cubic anomaly tensor is

$$d_R^{abc} = \text{Tr}_R(T^a \{T^b, T^c\}), \quad d_{\bar{R}}^{abc} = -d_R^{abc}. \quad (2.7)$$

Full hypers cancel between R and $\bar{R}$; self-conjugate real or pseudoreal representations have zero cubic tensor. The adjoint gauginos are in a real representation. Thus local anomaly cancellation follows structurally for the allowed matter schema; the checker need not tabulate every cubic invariant. For semisimple product groups, mixed tensors with one generator from another factor vanish by $\text{Tr } T^a = 0$. Source: Bilal, discussion of chiral fermions, conjugate representations and anomaly traces [5]; matter restrictions in [3].

The mod-two test

For factors with the conventional four-dimensional Witten anomaly, the number of anomalous Weyl representations must be even. The code checks A_1 , all supported C_r , and $B_2 \simeq C_2$. In its normalization the parity is

$$\mathfrak{a}_a = \sum_{h \text{ half}} n_h [2T_a(R_{ha})] \prod_{b \neq a} \dim R_{hb} \pmod{2}. \quad (2.8)$$

Full hypers contribute an even number automatically. The integer entering this test is $2T$, not the full-hyper beta contribution with an additional factor of two. A nonzero parity makes the gauge theory inconsistent within this setup. Source: Witten [6]; Wang–Wen–Witten, §2.3 [7]; symplectic-group discussion in [3, 8].

For $SU(2)$ isospin j , a useful independent formula is

$$T_j = \frac{1}{3}j(j+1)(2j+1), \quad 2T_j = \frac{2}{3}j(j+1)(2j+1). \quad (2.9)$$

It is odd for $j = 2k + 1/2$, $k \in \mathbb{Z}_{\geq 0}$: $j = 1/2, 5/2, 9/2, \dots$. One fundamental half hyper is anomalous, two are safe. A single $SU(2)^3$ trifundamental half hyper gives four doublets to each factor and is safe under this conventional test, although it is not by itself conformal.

Not an all-background anomaly classification

The project assumes the ordinary spin-theory setting and simply connected groups. It does not implement the newer non-spin $SU(2)$ anomaly associated with isospin $4k + 3/2$, global quotients, discrete theta angles or general bordism anomalies. Flavor 't Hooft anomalies need not vanish and are not gauge-anomaly rejection criteria. See [7, 8] for the distinction.

2.4 Independent $SU(N)$ formula checks

This section retains the useful representation-theory cross-checks from the older reference. They are derived checks, not a second implementation backend. For Dynkin labels $(a_1, \dots, a_{N-1})$, define row lengths and box count by

$$\ell_i = \sum_{k=i}^{N-1} a_k \quad (i < N), \quad \ell_N = 0, \quad B = \sum_{i=1}^N \ell_i. \quad (2.10)$$

Specializing the Weyl dimension and highest-weight Casimir formulas gives

$$\dim R = \prod_{1 \leq i < j \leq N} \frac{\ell_i - \ell_j + j - i}{j - i}, \quad (2.11)$$

$$C_2(R) = \frac{1}{2} \left[NB + \sum_{i=1}^N \ell_i(\ell_i + 1 - 2i) - \frac{B^2}{N} \right], \quad (2.12)$$

$$T(R) = \frac{\dim R}{N^2 - 1} C_2(R). \quad (2.13)$$

Source: Direct specialization of (2.2) and (2.3); compare Sage characters [2] and representation tables in [4].

$SU(N)$ representation	Dimension	C_2	T
Fundamental	N	$(N^2 - 1)/(2N)$	$1/2$
Adjoint	$N^2 - 1$	N	N
Two-index symmetric	$N(N+1)/2$	$(N-1)(N+2)/N$	$(N+2)/2$
Two-index antisymmetric	$N(N-1)/2$	$(N+1)(N-2)/N$	$(N-2)/2$

One symmetric plus one antisymmetric full hyper therefore gives $2T_{\text{sym}} + 2T_{\text{asym}} = 2N$ and $b_0 = 0$. For $N = 2$ the antisymmetric is a singlet, with zero T .

Function-to-equation map

Routine	Mathematical object
<code>get_lie_algebra, validate_dynkin_labels</code>	Root data and (2.1)
<code>representation_dimension</code>	(2.2), checked by (2.11)
<code>quadratic_casimir, dynkin_index</code>	(2.3)
<code>conjugate_dynkin_labels, representation_reality</code>	(2.4)
Gauge-factor beta calculation	(2.5)
Gauge-anomaly checks	(2.7), (2.8)

3 Package index

3.1 Full-index convention and single-letter functions

MODULE `index/n2_theory_index.py`

In radial quantization the chosen right-handed supercharge restricts the trace to

$$\delta_R = E - 2j_2 - 2R + r = 0, \quad (3.1)$$

$$\mathcal{I}(t, y, u) = \text{Tr}_{\delta_R=0} (-1)^F t^{2(E+j_2)} y^{2j_1} u^{-2(r+R)}. \quad (3.2)$$

This is the literature convention with $v = u^2$. It is a signed protected trace, not a positive count of every short multiplet: cancellations and recombination mean the full operator spectrum cannot generally be recovered uniquely. Source: Gadde–Pomoni–Rastelli, Eqs. (5.4)–(5.5) [9].

Writing $D(t, y) = (1 - t^3 y)(1 - t^3 / y)$, the scalar coefficients of the vector and half-hyper characters are

$$f_V(t, y, u) = \frac{t^2 u^2 - t^4 u^{-2} - t^3(y + y^{-1}) + 2t^6}{D(t, y)}, \quad (3.3)$$

$$f_H(t, y, u) = \frac{t^2/u - t^4 u}{D(t, y)}. \quad (3.4)$$

The denominator counts the two allowed spacetime derivatives; fermionic letters and equations of motion produce the numerator signs. Source: Gadde–Pomoni–Rastelli, Eqs. (5.21)–(5.22), with $v = u^2$ [9].

For gauge variables $g = (g_1, \dots, g_s)$, the complete letter function is

$$L = f_V \sum_a \chi_{\text{adj}_a}(g_a) + f_H \left[\sum_{h \text{ full}} n_h (\chi_{\mathcal{R}_h} + \chi_{\overline{\mathcal{R}_h}}) + \sum_{h \text{ half}} n_h \chi_{\mathcal{R}_h} \right]. \quad (3.5)$$

Here $\chi_{\mathcal{R}_h}(g) = \prod_a \chi_{\mathcal{R}_{ha}}(g_a)$. The conjugate of a product representation conjugates every factor simultaneously: $(R, S) \mapsto (\overline{R}, \overline{S})$. A full bifundamental is not four independent choices of conjugation. For a self-conjugate full hyper, the two equal characters still give multiplicity two.

Flavor-unrefined does not mean flavor-singlet

All flavor fugacities are currently set to one. Flavor representations are replaced by their dimensions, while gauge representations are kept until singlet projection. Charged flavor operators are counted; the code does not project them out. Spin and R-symmetry fugacities y, u are still present.

3.2 Plethystic expansion and FORM truncation

For a letter function with no degree-zero term, define

$$\text{PE}[L] = \exp \left[\sum_{k \geq 1} \frac{1}{k} L(t^k, y^k, u^k; g_1^k, \dots, g_s^k) \right], \quad \mathcal{I} = \int_G d\mu(g) \text{PE}[L]. \quad (3.6)$$

Adams substitution acts on every fugacity and character, but leaves numerical multiplicities and the signed letter coefficients as coefficients. The Haar measure is normalized so $\int_G 1 = 1$. Source: Plethystics [10]; gauge-index construction in [9].

`_character_basis` assigns formal character functions $C_p(k) = \chi_{R_p}(g_a^k)$ to simple-factor irreps. `_build_form_program` constructs the truncated exponent

$$A_D = \left[\sum_{k=1}^{\lfloor D/2 \rfloor} \frac{L(t^k, y^k, u^k; g^k)}{k} \right]_{t \leq D}. \quad (3.7)$$

Every letter begins at t^2 , so $k > \lfloor D/2 \rfloor$ and particle number $m > \lfloor D/2 \rfloor$ cannot contribute. Derivatives are expanded as

$$D(t, y)^{-1} = \sum_{p, q \geq 0} t^{3(p+q)} y^{p-q}; \quad p, q \leq \lfloor D/3 \rfloor \quad (3.8)$$

is a sufficient rectangular bound before final degree truncation. FORM declares t with maximum power D , so excess combined powers are discarded throughout the calculation. Negative powers of u and y are allowed.

The auxiliary-variable recurrence

The code starts with zA_D and successively substitutes

$$z \mapsto 1 + \frac{zA_D}{m}, \quad m = 2, \dots, M, \quad M = \lfloor D/2 \rfloor. \quad (3.9)$$

After step m , the expression is $\sum_{j=1}^{m-1} A_D^j/j! + zA_D^m/m!$. Setting $z = 1$ and adding the vacuum term gives $\sum_{j=0}^M A_D^j/j!$ exactly through degree D . This is the algebraic reason for the generated FORM substitutions, rather than an approximation requiring a convergence assumption. Source: Recurrence derived from the source; FORM syntax and truncation conventions in the supplied FORM 5 manual [12], with software background in [13].

`run_form` is imported from `common/form_utils.py`. `_parse_form_output` converts rational coefficients, fugacity exponents and character powers into `FormTerm` records. The numerical coefficients remain exact rationals, including cancellations generated by the index.

3.3 Gauge singlets, LiE and the reusable character cache

MODULE `index/char_decomposition_cache.py`

The Adams operation is

$$\psi^k \chi_R(g) = \chi_R(g^k), \quad \prod_{k \geq 1} (\psi^k \chi_R)^{n_k} = \sum_{\lambda} c_{\lambda} \chi_{\lambda}. \quad (3.10)$$

The right-hand side may be a *virtual* character: c_{λ} need not be nonnegative. For $SU(2)$, $\psi^2 \chi_2 = \chi_3 - \chi_1$, while $\chi_2^2 = \chi_3 + \chi_1$. Replacing the former with the latter would change the PE and corrupt gauge projection. Source: LiE manual, character operations [11]; Sage character operations [2].

`CharacterDecompositionCache` uses LiE Adams and tensor operations to obtain the coefficients in (3.10). Requests are keyed by Cartan type, Dynkin labels and the complete Adams-power tuple $(n_1, n_2, \dots)$. The weighted order

$$w = \sum_{k \geq 1} k n_k \quad (3.11)$$

groups entries into files; it is not alone a unique cache key. For example, `A1/A1_dynkin_1_adams_order_2.json` can contain different tuples of weight two. Dependency reuse, in-memory lookups and persistent JSON avoid repeating LiE work. Batch requests can use worker processes; locks and atomic replacement preserve entries sharing a cache file.

For each gauge factor, tensor products of the resulting simple characters are decomposed and only the trivial-representation coefficient is retained:

$$\int_{G_a} \chi_{\lambda}(g_a) d\mu_a = \delta_{\lambda,0}, \quad \text{mult}_{1,G}(\boxtimes_a V_a) = \prod_a \text{mult}_{1,G_a}(V_a). \quad (3.12)$$

`_project_terms` performs this factorwise calculation without expanding all gauge torus monomials in FORM. B_r, D_r continue to mean Spin representations, not a restriction to representations descending to SO quotients.

Public result and a regression example

`calculate_index(data, order, ...)` returns an exact flat Sage Laurent polynomial in (t, y, u) over $\mathbb{Q}$. `calculate_index_internal` accepts already parsed factors and hypers; `calculate_index_from_file` reads JSON, and `parse_index_polynomial` reverses the serialized polynomial format. For $SU(3)$ with six full fundamentals, the tested expansion is

$$\begin{aligned} \mathcal{I} = & 1 + (36u^{-2} + u^4)t^4 - u^2(y + y^{-1})t^5 \\ & + (40u^{-3} - 36 + u^6)t^6 + O(t^7). \end{aligned} \quad (3.13)$$

The displayed $O(t^7)$ is explanatory: the returned Laurent polynomial does not store a precision bound. Cache reuse improves repeated calculations, but the number of terms and tensor decompositions still grows rapidly with the cutoff.

3.4 Coulomb branch: invariant degrees and operator counting

MODULE `index/n2_theory_coulomb_branches.py`

The old module name `n2_theory_branches.py` has been replaced by the name above. A scalar Coulomb primary obeys $j_1 = j_2 = R = 0$ and $E = -r = \Delta$ in (3.1). Its weight is $(t^2 u^2)^\Delta$. For the ordinary scalar Coulomb chiral ring, let

$$\mathcal{I}_C(x) = \sum_{\Delta \geq 0} \dim(\mathcal{R}_C)_\Delta x^\Delta, \quad x = t^2 u^2. \quad (3.14)$$

For a Lagrangian semisimple gauge theory, the adjoint scalar ϕ has dimension one and the invariant ring is

$$\mathbb{C}[\mathfrak{g}_\mathbb{C}]^{G_c} \simeq \mathbb{C}[\mathfrak{h}_\mathbb{C}]^W = \mathbb{C}[P_1, \dots, P_\ell], \quad \Delta(P_i) = d_i = m_i + 1. \quad (3.15)$$

Here $\ell = \text{rank } G$, d_i are Weyl invariant degrees and m_i are Weyl exponents. Hence

$$\mathcal{I}_C(x) = \text{PE} \left[\sum_i x^{d_i} \right] = \prod_i \frac{1}{1 - x^{d_i}}. \quad (3.16)$$

Source: Invariant-degree interpretation and tables in Sage [14]; Casimir operators and Coulomb dimensions in Shapere–Tachikawa, §3 and §5.1 [15]; Coulomb-index letters in [16], Eq. (4.17).

Type	Basic invariant degrees (with multiplicity)
A_r	$2, 3, \dots, r + 1$
B_r, C_r	$2, 4, \dots, 2r$
D_r	$2, 4, \dots, 2r - 2$ together with r
E_6	$2, 5, 6, 8, 9, 12$
E_7	$2, 6, 8, 10, 12, 14, 18$
E_8	$2, 8, 12, 14, 18, 20, 24, 30$
F_4	$2, 6, 8, 12$
G_2	$2, 6$

`_invariant_degrees` caches `WeylGroup(cartan_type).degrees()`. `coulomb_branch_spectrum_from_gauge_factors` accepts one `GaugeFactorData` or an iterable and returns the sorted integer tuple for all factors. Multiplicities are preserved: D_4 gives $(2, 4, 4, 6)$, not $(2, 4, 6)$. Matter information is not needed for this invariant-ring computation. The group-only function does not independently test whether matter makes the theory conformal.

3.5 Coulomb PE with rational dimensions and FORM

The spectrum interface also accepts positive exact rational dimensions, so it can serve as a counting backend for future non-Lagrangian inputs. If $f(x) = \sum_{\Delta} c_{\Delta} x^{\Delta}$, then

$$\text{PE}[f] = \exp\left(\sum_{k \geq 1} \frac{f(x^k)}{k}\right) = \prod_{\Delta} (1 - x^{\Delta})^{-c_{\Delta}}. \quad (3.17)$$

The spectrum API requires $c_{\Delta} \in \mathbb{Z}_{\geq 0}$; the lower-level `calculate_plethystic_exponential` accepts signed integer coefficients. Negative entries give numerator factors and can describe complete-intersection relations, but an arbitrary signed list is not automatically a valid ring. Source: Plethystic exponential and product identity, [10], §5.

Clearing denominators exactly

For inclusive physical cutoff D , the implementation sets

$$L = \text{lcm}(\text{den}(D), \{\text{den}(\Delta) : c_{\Delta} \neq 0, \Delta \leq D\}), \quad q = x^{1/L}, \quad M = LD. \quad (3.18)$$

Every input monomial becomes $q^{L\Delta}$, with integer exponent. If the lowest nonzero exponent is ν , the exact finite computation is

$$A_M(q) = \sum_{k=1}^{\lfloor M/\nu \rfloor} \frac{1}{k} \sum_{\Delta} c_{\Delta} q^{kL\Delta}, \quad [\text{PE}(f)]_{q \leq M} = \left[\sum_{j=0}^{\lfloor M/\nu \rfloor} \frac{A_M(q)^j}{j!} \right]_{q \leq M}. \quad (3.19)$$

`_build_plethystic_form_program` uses FORM's bounded q symbol and the exponential recurrence (3.9). `_parse_form_series` reads exact rational coefficients; `_to_coulomb_series` maps q^m back to $x^{m/L}$ in `PuiseuxSeriesRing(QQ, "x")`. Empty retained input returns one without starting FORM. The scaling formulas are derived directly from the source; no floating-point fit to exponents is used.

Two useful examples

For $SU(3)$, independent generators have dimensions two and three:

$$\begin{aligned} \mathcal{I}_C &= (1 - x^2)^{-1} (1 - x^3)^{-1} \\ &= 1 + x^2 + x^3 + x^4 + x^5 + 2x^6 + x^7 + 2x^8 + O(x^9). \end{aligned} \quad (3.20)$$

The two terms at degree six are products P_2^3 and P_3^2 ; neither is a new generator. With one input dimension $6/5$ and cutoff $18/5$, the result is $1 + x^{6/5} + x^{12/5} + x^{18/5}$.

Mathematical input is not physical classification

Positive rational dimensions are accepted algebraically, including values that would violate the scalar unitarity bound in a unitary four-dimensional SCFT. The code does not establish that a supplied non-Lagrangian spectrum is realized by a theory or that its chiral ring is freely generated.

3.6 Taking the Coulomb limit of an existing full index

`calculate_coulomb_branch_index_from_full_index` accepts the project's Sage Laurent polynomial or its serialized string. Substitute

$$t \longrightarrow 0, \quad u = \frac{x^{1/2}}{t}, \quad y \text{ fixed}; \quad t^a y^b u^c = t^{a-c} y^b x^{c/2}. \quad (3.21)$$

Thus a canonical nonzero monomial has the following disposition:

Condition	Implementation
$a > c$	Vanishes; discard it.
$a = c, b = 0, c \geq 0$	Keep its coefficient at dimension $\Delta = c/2$.
$a < c$	Raise a divergent-limit error.
$a = c, b \neq 0$	Reject a surviving spin-dependent term.
$a = c, c < 0$	Reject a negative Coulomb dimension.

An optional `max_dimension` further truncates the retained terms. Like terms are already combined in the source Laurent ring. This path is coefficient selection in Sage/Python and does not run FORM or LiE; symbolic PE/PL is where FORM is used. Source: Project-variable derivation from (3.2); in the same limit (3.3) and (3.4) become $f_V = x, f_H = 0$, matching [16], Eq. (4.17).

Precision cannot be created by taking a limit

If the original index is known through t^N (inclusive), then x^Δ comes from $t^{2\Delta} u^{2\Delta}$. Guaranteed Coulomb information therefore extends only to

$$\Delta \leq N/2. \quad (3.22)$$

An order-18 full index can determine Coulomb coefficients through dimension nine, not ninety. The property workflow obtains the order-90 Coulomb index directly from Weyl degrees instead. From (3.13), the supported Coulomb result is $1 + x^2 + x^3$ through dimension three.

Caller-owned cutoff and adapter restrictions

The full-index ring and the returned finite Coulomb expressions do not carry the provenance of a truncated polynomial. The caller must retain the trusted cutoff and must not interpret omitted coefficients above it as zero. The full index adapter accepts integer Laurent exponents in t, y, u ; it is not a generic adapter for arbitrary fractional-power non-Lagrangian full indices. Direct Coulomb spectrum/series input supports rational dimensions more generally.

For non-Lagrangian applications this adapter also assumes the desired limit is the ordinary scalar Coulomb ring. More general protected contributions are not identified or decomposed by the code. Rejection of spin-dependent terms is a guard, not a proof of all the required physical assumptions.

3.7 Plethystic logarithm: recovering generator dimensions

For $H(0) = 1$, the inverse plethystic transform is

$$\text{PL}[H](x) = \sum_{k \geq 1} \frac{\mu(k)}{k} \log H(x^k), \quad \text{PL}[\text{PE}(f)] = f. \quad (3.23)$$

$\mu(1) = 1$; $\mu(k) = 0$ when k contains a squared prime factor. Otherwise $\mu(k) = (-1)^s$, where s is the number of distinct prime factors of k . Source: Plethystic logarithm/Möbius inversion, [10], §5.

`calculate_plethystic_logarithm` normalizes the input to a finite exact series, requires its constant coefficient to be one, rejects negative powers, checks any available finite-series precision, and rescales by (3.18). For input $H + O(x^P)$, the inclusive cutoff must satisfy $D < P$; a plain polynomial cannot enforce this provenance check. Write the resulting input as $1 + \delta(q)$ with lowest nonzero power ν and $K = \lfloor M/\nu \rfloor$. FORM evaluates

$$[\log(1 + \delta)]_{q \leq M} = \left[\sum_{j=1}^K \frac{(-1)^{j+1}}{j} \delta^j \right]_{q \leq M}. \quad (3.24)$$

`_build_plethystic_log_form_program` starts with $z\delta$ and uses

$$z \mapsto 1 - \frac{j-1}{j} z\delta, \quad j = 2, \dots, K. \quad (3.25)$$

After step j , this gives the completed terms up to δ^{j-1} plus $z(-1)^{j+1}\delta^j/j$; setting $z = 1$ proves the recurrence.

If FORM returns $\log H(q) = \sum_m b_m q^m$, Python then applies the exact Möbius coefficient transform

$$[q^n] \text{PL}(H) = \sum_{k|n} \frac{\mu(k)}{k} b_{n/k}. \quad (3.26)$$

The helper names are `_mobius` and `_apply_mobius_transform`. The returned Puiseux series contains signed rational coefficients; FORM performs the ordinary-log polynomial algebra, while Python handles this arithmetic transform.

What the extraction wrapper guarantees

`extract_coulomb_branch_spectrum_from_index` expands the PL and returns a sorted tuple containing each dimension as many times as its coefficient. It requires every retained coefficient to be a nonnegative integer; otherwise it raises an error. Use the raw PL routine when signed data is required. For (3.20), the PL through dimension eight is $x^2 + x^3$ and the tuple is (2, 3). For $(1 - x^4)/(1 - x^2)^2$, the PL is $2x^2 - x^4$; the raw result retains the relation and the free-spectrum wrapper rejects it.

A finite nonnegative PL is not a global freeness proof

Generators above the cutoff and relations first appearing later are invisible. For non-complete-intersection rings, higher PL terms can encode syzygies and their cancellations; positive terms are not universally all minimal generators. The wrapper produces a free-generator interpretation only through the supplied trusted cutoff, under the stated ring assumptions.

3.8 Coulomb API guide and reuse boundaries

All names below are in `index/n2_theory_coulomb_branches.py`.

Entry point	Input and result
<code>coulomb_branch_spectrum_from_gauge_factors</code>	One/iterable of <code>GaugeFactorData</code> ; integer dimension tuple. Uses Sage degrees only.
<code>calculate_lagrangian_coulomb_branch_index</code>	Gauge-factor data and cutoff; Weyl spectrum $\rightarrow$ FORM PE $\rightarrow$ exact x series.
<code>calculate_coulomb_branch_index</code>	Exactly one of <code>spectrum</code> or <code>full_index</code> . Spectrum requires a cutoff; full-index mode permits it.
<code>calculate_coulomb_branch_index_from_full_index</code>	Laurent polynomial/string; select terms by (3.21).
<code>calculate_plethystic_exponential</code>	Positive rational dimensions mapped to signed integer coefficients, plus cutoff.
<code>calculate_plethystic_logarithm</code>	Unit-constant Coulomb series/string and trusted cutoff; signed rational Puiseux series.
<code>extract_coulomb_branch_spectrum_from_index</code>	Same index input; repeated <code>Fraction</code> dimensions if the retained PL is non-negative integral.
<code>parse_coulomb_branch_index</code>	Arithmetic-only serialized x expression; exact Sage Puiseux expression.

Minimal example, run with Sage Python

```
from fractions import Fraction
from anomalies.check_n2_anomalies import GaugeFactorData
from anomalies.lie_algebra import get_lie_algebra
from index import n2_theory_coulomb_branches as cb

gauge = GaugeFactorData("g", get_lie_algebra("A2"))
H = cb.calculate_lagrangian_coulomb_branch_index(gauge, 8)
dims = cb.extract_coulomb_branch_spectrum_from_index(H, 8)
assert dims == (Fraction(2), Fraction(3))

H_fractional = cb.calculate_coulomb_branch_index(
    [Fraction(6, 5)], Fraction(18, 5)
)
```

String parsers permit only the documented arithmetic variables and syntax, not general Sage code. A textual $O(x^D)$ precision annotation is not the serialized polynomial format. Supply exact fractions, not float approximations such as 1.2, and retain the cutoff separately when serializing a truncated result.

The common property module has similarly named functions taking a *theory dictionary*, not `GaugeFactorData`. Use explicit imports to distinguish the two API layers. Its Lagrangian spectrum path reads Weyl degrees directly; it does not need to compute a PL of its already-known spectrum.

3.9 Future extension: Hall–Littlewood and Higgs operators

Mathematical design notes only

There is currently no Higgs/HL calculation module, flavor-refined index output, or Higgs-spectrum database property. This section records the mathematical background and proposed implementation changes from the session, not existing functionality.

Higgs-branch chiral primaries are scalar $\widehat{\mathcal{B}}_R$ primaries with

$$E = 2R, \quad r = 0, \quad j_1 = j_2 = 0. \quad (3.27)$$

The HL trace instead imposes $\delta_R = 0$ and $\delta_{1\pm} = E \pm 2j_1 - 2R - r = 0$. Solving these equations gives

$$j_1 = 0, \quad j_2 = r, \quad E = 2R + r, \quad \mathcal{I}_{\text{HL}} = \text{Tr}_{\text{HL}}(-1)^F \tau^{2(E-R)}. \quad (3.28)$$

These are conditions on contributing states, including possible fermionic cohomology representatives, not only Higgs scalar primaries. On (3.27), the exponent becomes $E = 2R$. This explains why the two sets of conditions differ without contradicting one another. Source: Gadde–Rastelli–Razamat–Yan, Eqs. (4.1), (4.5)–(4.8) and Appendix B [16]; general caveats about identifying HL with a Higgs Hilbert series in [19].

The limit in the project's variables

Keeping $\tau = t^2/u$ fixed and sending $t \rightarrow 0$ gives

$$u = t^2/\tau, \quad t^a y^b u^c = t^{a+2c} y^b \tau^{-c}, \quad f_H \rightarrow \tau, \quad f_V \rightarrow -\tau^2. \quad (3.29)$$

Thus $a + 2c = 0$ selects surviving terms, positive values vanish, and negative values signal a divergent limit. Surviving y dependence should cancel for an HL result. A full index known to t^N supports integer HL degree at most $\lfloor N/2 \rfloor$. Direct HL construction would instead use the degree-one hyper letter, so its Adams/particle bound is D , not the current full-index $\lfloor D/2 \rfloor$.

Why the vector letter represents quadratic constraints

For full hypers, the $\mathcal{N} = 2$ superpotential can be written

$$W = \sqrt{2} \sum_a \Phi_a^A \mu_a^A(Q, \tilde{Q}), \quad \mu_a^A = \sum_h \tilde{Q}_h T_{R_{ha}}^A Q_h. \quad (3.30)$$

At $\Phi = 0$, the nontrivial F-term equations are $\partial W / \partial \Phi_a^A = \sqrt{2} \mu_a^A = 0$. They are quadratic in hyper scalars and transform in adj_a ; derivatives with respect to the hypers vanish there. Half-hyper notation uses the invariant symplectic pairing and gives the same quadratic moment-map structure. Source: Superpotential/F-term argument in [16], Eqs. (4.9)–(4.11); general Lagrangian background in [17].

3.10 Future extension: constraints, syzygies and flavor refinement

Let $S = \mathbb{C}[Q, \tilde{Q}]$ denote the polynomial ring of complex hyper coordinates. The Higgs chiral ring in the chosen complex structure is

$$\mathcal{R}_H = (S/(\mu_a^A))^{G_{\mathbb{C}}}. \quad (3.31)$$

If the moment-map components form a homogeneous regular sequence, quotienting by them multiplies the equivariant Hilbert series by

$$\text{PE} \left[-\tau^2 \sum_a \chi_{\text{adj}_a}(g_a) \right], \quad H_H(\tau) = \int_G d\mu \text{PE} \left[\tau \chi_{\text{hyp}} - \tau^2 \chi_{\text{adj}} \right]. \quad (3.32)$$

For one homogeneous non-zero-divisor f of degree d , the exact sequence $0 \rightarrow S(-d) \xrightarrow{\times f} S \rightarrow S/(f) \rightarrow 0$ gives the factor $1 - \tau^d$. Iterating is valid for a regular sequence. More generally, the fermionic factor is the Euler characteristic of the Koszul complex with odd variables η_a^A and differential $d\eta_a^A = \mu_a^A$. Its zeroth homology is the quotient; higher homology can obstruct equality with the HL index. Source: Regular/Koszul sequences [20]; physical HL/Higgs caveat [19].

For $SU(2)$, explicitly,

$$\text{PE}[-\tau^2 \chi_{\mathbf{3}}(a)] = (1 - \tau^2 a^2)(1 - \tau^2)(1 - \tau^2 a^{-2}) = 1 - \tau^2 \chi_{\mathbf{3}} + \tau^4 \chi_{\mathbf{3}} - \tau^6. \quad (3.33)$$

The higher terms are the alternating exterior powers of the constraint space. One must check assumptions for the theory in question; a blanket statement that every HL index, or every genus-zero construction, equals a Higgs Hilbert series is too strong [19].

What flavor-unrefinement loses

A refined Hilbert series organizes graded operators as

$$H_H(\tau, \mathbf{a}) = \sum_{\Delta, \lambda} m_{\Delta, \lambda} \chi_{\lambda}^F(\mathbf{a}) \tau^{\Delta}, \quad H_H(\tau, \mathbf{1}) = \sum_{\Delta} \underbrace{\sum_{\lambda} m_{\Delta, \lambda} \dim \lambda}_{N_{\Delta}} \tau^{\Delta}. \quad (3.34)$$

Unrefined data retains graded dimensions N_{Δ} but cannot uniquely recover flavor representations. It is sufficient if only those counts or an unrefined PL are desired, subject to the same Higgs/HL and syzygy caveats. Highest-weight generating functions are a possible later representation-level summary [21].

Proposed implementation changes

Preserve flavor characters in the letter builder and under every Adams operation. Complex blocks use $R \otimes \mathbf{n}$ and $\bar{R} \otimes \bar{\mathbf{n}}$; real blocks use the $Sp(n)$ defining representation; pseudoreal blocks use the $SO(m)$ vector. Project only gauge factors, leaving flavor characters in the output. Simple-factor LiE/cache machinery can be reused, but the coefficient ring, parsers and serialization must support flavor labels and explicit abelian flavor charges. A refined PL must substitute *all* flavor variables. Share flavor-block construction in a lower-level helper to avoid a circular import between properties and index.

4 Package common

4.1 Flavor symmetry and exactly marginal gauge couplings

MODULE `common/n2_theory_properties.py`

The property layer reuses anomaly-validated matter and groups it into flavor blocks. The symmetry is the commutant of the gauge action on the hyper space, with its compatible reality structure. At zero masses the connected factors are

$$F_{\mathcal{R}} = \begin{cases} U(n), & \mathcal{R} \text{ complex, } n \text{ full hypers,} \\ Sp(n), & \mathcal{R} \text{ real, } n \text{ full hypers,} \\ SO(m), & \mathcal{R} \text{ pseudoreal, } m = 2n_{\text{full}} + n_{\text{half}}. \end{cases} \quad (4.1)$$

Here $Sp(n)$ has rank n and defining dimension $2n$, often written $USp(2n)$. Complex-conjugate labels are merged into a single block, including simultaneous conjugation of all product factors. A pseudoreal block counts half-hyper units. An $SO(1)$ block has no continuous flavor symmetry. Source: Hyper reality in [3]; flavor/commutant background in [17]; the explicit grouping rule is the project's implementation of that construction.

The helpers `_simple_flavor_blocks`, `_product_flavor_blocks`, `_canonical_complex_key` and `_calculate_flavor_symmetry` carry out this grouping. This is flavor-symmetry metadata only: it is not used to refine the current index. Disconnected components, faithful global quotients and strong-coupling accidental enhancements are not determined by this routine.

For an anomaly-free conformal Lagrangian theory with s simple factors, the code assigns one exactly marginal complex gauge coupling per factor:

$$\tau_a = \frac{\theta_a}{2\pi} + \frac{4\pi i}{g_a^2}, \quad \dim_{\mathbb{C}} \mathcal{M}_{\text{conf}} = s. \quad (4.2)$$

This is the local dimension of the Lagrangian conformal manifold, not a description of its global S-duality quotient or singular loci. The τ_a in this equation is a coupling, distinct from the HL grading τ . Source: Lagrangian conformality [3]; $\mathcal{N} = 2$ supersymmetry-preserving marginal deformations in Córdova–Dumitrescu–Intriligator, §3.2.2/Table 22 [18].

Examples

$SU(3)$ with six full fundamental hypers has a $U(6)$ flavor block and one exactly marginal gauge coupling. $SU(2)$ with four full fundamentals is instead a pseudoreal block with eight half-hyper units, hence $SO(8)$ flavor. The distinction follows from total representation reality, not just multiplicity.

4.2 Central charges and property orchestration

For a conformal Lagrangian realization, define the effective numbers of vector and hypermultiplets by representation dimensions:

$$n_v = \sum_a \dim \mathfrak{g}_a, \quad n_H = \sum_h n_h \epsilon_h \prod_a \dim R_{ha}. \quad (4.3)$$

The Weyl-anomaly coefficients are

$$a = \frac{5n_v + n_H}{24}, \quad c = \frac{2n_v + n_H}{12}. \quad (4.4)$$

These include free gauge-singlet hypers. The code checks positivity and the equivalent identities $4(2a - c) = n_v$, $4(5c - 4a) = n_H$. Source: Shapere–Tachikawa, free-multiplet values in Eqs. (2.5)–(2.6) [15].

An additional mathematical cross-check with the Lagrangian Coulomb spectrum is

$$4(2a - c) = \sum_i (2d_i - 1) = \dim \mathfrak{g}. \quad (4.5)$$

The final equality follows from $\sum_i (d_i - 1) = |\Phi^+|$ and $\dim \mathfrak{g} = \ell + 2|\Phi^+|$. The more general Coulomb-dimension relation requires the assumptions discussed in [15], §5.1, Eq. (5.4); the project does not infer arbitrary non-Lagrangian central charges from it. The cross-check with degrees is explanatory here, not a separate implemented validation step. For $SU(3)$ with six full fundamentals, $n_v = 8$, $n_H = 18$, so $a = 29/12$, $c = 17/6$ and $4(2a - c) = 3 + 5 = 8$.

Current property pipeline

`calculate_n2_theory_properties(data)` validates the theory once and uses the internal gauge/hyper data for subsequent work. If it is an SCFT candidate:

1. Compute flavor metadata, gauge marginal couplings and exact central charges.
2. Compute the full index through `INDEX_MAX_ORDER=18`.
3. Build the Coulomb index directly from gauge degrees through `C_INDEX_MAX_ORDER=90`, independent of the full-index cutoff.
4. Return Coulomb generator dimensions directly from the same degrees.

For valid non-candidates, SCFT-dependent values are `None`; flavor metadata is still available. Malformed input raises an error.

Property key	Representation in aggregate result
<code>central_charges</code>	Exact rational a, c only.
<code>superconformal_index</code>	Serialized Laurent polynomial string.
<code>coulomb_branch_index</code>	Serialized x -series expression string.
<code>coulomb_branch_spectrum</code>	Tuple of exact <code>Fraction</code> dimensions.
<code>flavor_symmetry</code>	Flavor-factor metadata.
<code>conformal_manifold_dimension</code>	Number of marginal gauge couplings.

The public single-property wrappers return the corresponding exact Sage object or tuple before aggregate string serialization. The additional keys include `group`, `lagrangian_scft_candidate` and `exactly_marginal_gauge_couplings`.

4.3 Shared exact arithmetic, FORM execution and JSON

MODULE `common/number_utils.py`; `common/form_utils.py`; `common/json_utils.py`

These modules hold common utilities formerly duplicated in physics modules. Their role is to keep exact arithmetic and external-process handling consistent.

Exact-number contracts

Routine	Validation
<code>as_integer</code>	Rejects Python booleans/floats, including 1.0; requires exact equality with the converted integer.
<code>as_nonnegative_int</code>	Calls <code>as_integer</code> , then requires $n \geq 0$.
<code>as_nonnegative_fraction</code>	Accepts supported exact rational objects or integer/fraction text; rejects Python booleans/floats and negative values.
<code>as_positive_fraction</code>	Adds the strict $q > 0$ condition.

Use `Fraction(6,5)` or Sage `QQ(6)/5` for a Coulomb dimension; use integer counts for Dynkin labels, hyper multiplicities and cutoffs in t . The rational-text format is an integer or fraction such as "6/5", not decimal text such as "1.2". Accepting integral floats would broaden the shared API everywhere, not just in the Coulomb module; the current behavior is deliberately strict. These validators are type/value checks, not physical unitarity or representation-existence proofs.

One FORM execution boundary

`run_form(program, form_executable=..., timeout=...)` writes a FORM program in an isolated temporary directory, invokes the executable with quiet output, and returns the output text. It turns missing executables, timeouts, nonzero exit status or reported stderr into errors with useful context. Both the full index and Coulomb PE/PL reuse this function; there is no need for separate subprocess wrappers in each module.

`split_top_level` and `split_signed_terms` support parsing expressions without splitting inside parentheses or mistaking an exponent sign for a new term. Physics-specific parsers remain with their own modules because character factors and Puiseux powers require different validation. Source: Execution and parser contracts from the source; FORM program syntax in [12].

Serialization is exact, but precision metadata is separate

The JSON helpers encode a fraction as

```
{"numerator": 6, "denominator": 5}
```

Tuples become arrays, dataclass records become structured objects, and the aggregate property layer serializes Sage indices as expression strings. For example, the spectrum (2,3) is stored as an array of the two exact fractions, not as the polynomial $x^2 + x^3$.


Current index strings do not embed their trusted truncation order. Future non-Lagrangian imports should carry cutoff/provenance alongside the expression before any PL extraction or equality comparison; no such schema extension is claimed to exist in this edition.

4.4 MySQL persistence and schema version 6

MODULE `common/n2_theory_db.py`

The database distinguishes a theory, its shared properties, and its realizations. Flavor factors belong to the theory; gauge factors, hypermultiplets and exactly marginal couplings belong to a Lagrangian realization. The non-Lagrangian table is a placeholder, with no implemented import workflow.

Database relationships



Arrows run from parent to child; labels give the number of children per parent allowed by the schema. Each child has exactly one parent along each foreign key. PK = primary key; FK = foreign key; UQ = unique key. “Realization ID” abbreviates `lagrangian_realization_id`. All declared FKs use `ON DELETE CASCADE`.

How product-group matter is represented

One `hypermultiplets` row holds a complete tensor-product representation and its submitted `kind` and multiplicity. Its rows in `hypermultiplet_representations` identify the representation under each gauge factor, including singlets. The latter table joins matter to gauge factors; it does not create a separate hypermultiplet for each factor.

Shared flavor versus the submitted matter split

Schema 6 retains `half_hyper_units` in the shared flavor table and JSON, with $m = 2n_{\text{full}} + n_{\text{half}}$. The original split remains in realization rows and input/anomaly JSON. Public property-calculator output still reports both counts; the database strips them from its shared-property copy.

4.5 Database table catalog and keys

All tables use InnoDB and utf8mb4. IDs are unsigned BIGINTs; standalone id keys auto-increment. Exact rational values use JSON numerator/denominator records or separate numeric columns, as detailed below.

Table	Keys and stored data
<code>schema_metadata</code>	PK <code>metadata_key</code> ; <code>metadata_value</code> holds the schema version (current value "6"). Independent of theory rows.
<code>theories</code>	PK <code>id</code> ; UQ <code>canonical_key</code> . Display name and creation/update timestamps. A new Lagrangian theory uses <code>lagrangian:<hash></code> as its key.
<code>theory_properties</code>	PK/FK <code>theory_id</code> $\rightarrow$ <code>theories.id</code> . Flavor group/rank/dimension, conformal-manifold dimension, exact central-charge JSON, separate full-index, Coulomb-index and Coulomb-spectrum JSON, and combined <code>properties_json</code> . Generated, indexed DECIMAL(65,30) columns approximate a, c for queries.
<code>flavor_symmetry_factors</code>	PK <code>id</code> ; FK <code>theory_id</code> $\rightarrow$ <code>theories.id</code> ; UQ (<code>theory_id</code> , <code>factor_order</code>). Group, Lie algebra/rank/dimension, representation reality, <code>half_hyper_units</code> , and <code>gauge_representation_json</code> . No full/half split columns in schema 6.
<code>lagrangian_realizations</code>	PK <code>id</code> ; FK <code>theory_id</code> $\rightarrow$ <code>theories.id</code> ; UQ <code>canonical_hash</code> (64-character SHA-256). Gauge group/count, anomaly and conformality flags, input/anomaly/coupling JSON, creation timestamp.
<code>non_lagrangian_realizations</code>	PK <code>id</code> ; FK <code>theory_id</code> $\rightarrow$ <code>theories.id</code> . Construction type, optional source reference, <code>data_json</code> , creation timestamp. Reserved.
<code>gauge_factors</code>	PK <code>id</code> ; FK realization ID $\rightarrow$ <code>lagrangian_realizations.id</code> . Separate UQs on (realization ID, <code>factor_order</code>) and (realization ID, <code>factor_key</code>). Group, Cartan type/rank, exact vector, matter and b_0 fractions, beta/global-anomaly flags and Witten parity.
<code>hypermultiplets</code>	PK <code>id</code> ; FK realization ID $\rightarrow$ <code>lagrangian_realizations.id</code> ; UQ (realization ID, <code>hyper_order</code>). Name, <code>kind</code> (full/half), submitted multiplicity, total tensor-product dimension and reality.
<code>hypermultiplet_representations</code>	Composite PK (<code>hypermultiplet_id</code> , <code>gauge_factor_id</code>); separate FKs to <code>hypermultiplets.id</code> and <code>gauge_factors.id</code> . Name, Dynkin-label JSON, dimension/reality, exact Dynkin-index and beta-contribution numerators/denominators.
<code>exactly_marginal_couplings</code>	Composite PK (realization ID, <code>coupling_order</code>); FK realization ID $\rightarrow$ <code>lagrangian_realizations.id</code> . Stores <code>gauge_factor_key</code> as text; that key has no additional FK to <code>gauge_factors</code> .

Serialization and integrity boundaries

Full and Coulomb indices are JSON expression strings. Coulomb spectra are arrays of exact fractions; central charges are exact a, c fraction records. The separate JSON fields are repeated in `properties_json`. Gauge beta fractions and representation invariants use DECIMAL(65,0) integer columns. The writer keeps both ends of each matter–gauge relation within the same realization; the two individual FKs do not themselves enforce that match.

4.6 Canonical matter, migrations and import behavior

The canonical Lagrangian hash is SHA-256 of a normalized payload. Names resolve to Dynkin labels, repeated matter entries combine, and zero multiplicities vanish. Conjugate choices use the lexicographically larger label tuple, favoring fundamentals. Product representations identify only simultaneous conjugation of all factors. Matter order is ignored; gauge-factor order remains significant.

Full/half equivalence for pseudoreal representations

For each complete pseudoreal gauge representation, aggregate before pairing:

$$m = 2n_{\text{full}} + n_{\text{half}}, \quad (n_{\text{full}}^{\text{canonical}}, n_{\text{half}}^{\text{canonical}}) = (\lfloor m/2 \rfloor, m \bmod 2).$$

Four full $SU(2)$ fundamentals, eight halves, or two full plus four halves now have the same hash. An odd remaining half stays in its own irrep. For products, the rule uses the reality of the entire tensor product. Real/complex matter remains full; invalid half-hyper input is rejected by the checker.

Schema history

Migration	Change
1 → 2	Add generated, indexed decimal central charges; exact JSON remains authoritative.
2 → 3	Rename the plural full-index property and column to singular <code>superconformal_index</code> / <code>superconformal_index_json</code> .
3 → 4	Rename the old Coulomb-spectrum/index placeholder to <code>coulomb_branch_index_json</code> , including its combined-JSON key.
4 → 5	Add nullable <code>coulomb_branch_spectrum_json</code> for actual generator dimensions; do not compute spectra for existing rows.
5 → 6	Drop <code>full_hypermultiplets</code> and <code>half_hypermultiplets</code> from <code>flavor_symmetry_factors</code> ; keep <code>half_hyper_units</code> and the original realization records.

Schema upgrade and data migration are different operations

Version 6 does not rehash existing realizations, merge duplicate theories, or eagerly rewrite old shared JSON. Old hashes from the earlier conjugacy or full/half conventions require a separate data migration. Under the current conjugacy convention, inputs already written entirely as full hypers keep their hashes after full/half normalization.

`_insert_shared_properties` ignores the obsolete flavor split when comparing legacy JSON. Matching JSON is rewritten without that split; a missing Coulomb spectrum can also be backfilled. Physical property mismatches still raise an error. Shared flavor metadata retains gauge labels and factor IDs, so this is not a general comparison of arbitrary dual realizations.

`store_lagrangian_theory` checks anomalies/conformality and calculates properties before looking up the hash. A duplicate returns the existing IDs with `inserted=False`, without updating the realization or calling the shared-property helper; it therefore does not trigger legacy-JSON cleanup or spectrum backfill. A new import uses parameterized SQL with commit/rollback. Schema initialization precedes that DML transaction; MySQL DDL is not covered by its rollback. No live database was changed for this documentation update.

5 Package test and reproducible use

Selected test modules	Main coverage
<code>test_check_n2_anomalies.py</code>	Root/representation data, all families, spectator factors, full/half validation, beta and Witten checks.
<code>test_n2_theory_index.py</code>	Laurent output, named/explicit labels, product singlets, low-order reference indices, LiE parsing and cache concurrency.
<code>test_n2_theory_coulomb_branches.py</code>	Weyl degrees, repeated/fractional dimensions, PE/PL round trips, relations, full-index projection and precision.
<code>test_n2_theory_properties.py</code>	Flavor blocks, central charges, property keys, direct gauge-based Coulomb wiring and exact JSON.
<code>test_n2_theory_db.py</code>	Schema/version handling, shared property comparison, full/half deduplication, legacy JSON/spectrum backfill, SQL parameters and transactions; actual index equality through t^6 .

5.1 Validation snapshot

The recorded 9 September 2026 run after the database fix completed in 9.507 s:

134 tests run: 132 passed, 2 skipped.

The two skipped tests require live MySQL. Database unit tests use recording/mock connections; a backend regression compares actual FORM/LiE indices for full/half $SU(2)$ fundamentals and $SU(2)^3$ trifundamentals. Other tests cover Tits' reality formula, irrep/theory enumeration and index/Coulomb calculations. This recorded run does not validate live migration or arbitrary high-order performance.

5.2 Commands and dependencies

Run from the repository root with Sage's Python, not an unrelated Python interpreter. FORM and LiE must be discoverable for backend tests.

```
sage -python -m unittest discover -s test -p 'test_*.py' -v
sage -python -m index.n2_theory_index --help
sage -python -m common.n2_theory_properties --help
sage -python -m common.n2_theory_db --help
```

Sage supplies exact rings and Lie data; FORM and LiE handle expansions and representation decompositions. PyMySQL supplies the database connection. The low-level APIs accept executable paths, timeouts and cache directories.

5.3 Practical boundaries before extending the project

Preserve the exact-number contracts, distinct index/spectrum properties and source precision. Non-Lagrangian counting needs explicit ring/freeness assumptions; Higgs work needs an HL/Hilbert-series justification; flavor refinement must preserve characters through gauge projection. These are mathematical contracts.

References and equation provenance

Equation numbers in this document are local. Citations near equations identify the primary source or software reference; algebraic recurrences, cutoff rules and function mappings are derived from the inspected project source. Literature section/equation numbers refer to the linked arXiv editions where specified. Software documentation was consulted on 8 September 2026. This bibliography also retains the relevant sources from both earlier summaries.

References

- [1] The Sage Developers, *Root systems*, SageMath reference manual. Dynkin labels, root/weight lattices and finite Cartan data. [Official root-system documentation](#).
- [2] The Sage Developers, *Weyl character rings*, SageMath reference manual. Character degrees, duals and Frobenius–Schur indicators. [Official Weyl-character documentation](#).
- [3] L. Bhardwaj and Y. Tachikawa, *Classification of 4d $\mathcal{N} = 2$ gauge theories*, JHEP 12 (2013) 100, updated arXiv edition. [arXiv:1309.5160](#). §2, Eqs. (2.1)–(2.8): matter, spectator dimensions and beta functions.
- [4] R. Feger, T. W. Kephart and R. J. Saskowski, *LieART 2.0 – A Mathematica Application for Lie Algebras and Representation Theory*, Comput. Phys. Commun. 257 (2020) 107490. [arXiv:1912.10969](#). Representation tables are cross-checks; LieART is not the project backend.
- [5] A. Bilal, *Lectures on Anomalies*. [arXiv:0802.0634](#). Chiral anomaly traces, conjugation and product-group anomaly structure.
- [6] E. Witten, *An $SU(2)$ Anomaly*, Phys. Lett. B 117 (1982) 324–328. [doi:10.1016/0370-2693\(82\)90728-6](#). Conventional four-dimensional mod-two gauge anomaly.
- [7] J. Wang, X.-G. Wen and E. Witten, *A New $SU(2)$ Anomaly*, J. Math. Phys. 60 (2019) 052301. [arXiv:1810.00844](#). §2.3 includes the conventional isospin parity formula and distinguishes the additional non-spin anomaly.
- [8] I. García-Etxebarria and M. Montero, *Dai–Freed anomalies in particle physics*, JHEP 08 (2019) 003. [arXiv:1808.00009](#). Global-anomaly scope beyond local anomaly polynomials.
- [9] A. Gadde, E. Pomoni and L. Rastelli, *The Veneziano Limit of $\mathcal{N} = 2$ Superconformal QCD: Towards the String Dual of $\mathcal{N} = 2$ $SU(N_c)$ SYM with $N_f = 2N_c$* , [arXiv:0912.4918](#). §5, Eqs. (5.4)–(5.5), (5.21)–(5.22): the right-index convention and letters.
- [10] S. Benvenuti, B. Feng, A. Hanany and Y.-H. He, *Counting BPS Operators in Gauge Theories: Quivers, Syzygies and Plethystics*, JHEP 11 (2007) 050. [arXiv:hep-th/0608050](#). §5: PE, inverse plethystics, complete intersections and refined counting.
- [11] M. A. A. van Leeuwen, A. M. Cohen and B. Lisser, *LiE: A Package for Lie Group Computations* (1992), and LiE reference manual, character operations (Adams and tensor products). [TU Eindhoven publication record](#); [CWI manual record](#).
- [12] J. A. M. Vermaseren, J. Davies, T. Kaneko, J. Kuipers, C. Marinissen, B. Ruijl, M. Tentyukov, T. Ueda and J. Vollinga, *FORM version 5.0.0 Reference manual*, 27 January 2026. User-supplied [form-5.0.0-manual.pdf](#); see Chapters 1–3 and the entries for symbol restrictions, substitution, sorting and rational functions.
- [13] J. Kuipers, T. Ueda, J. A. M. Vermaseren and J. Vollinga, *FORM version 4.0*, Comput. Phys. Commun. 184 (2013) 1453–1467. [arXiv:1203.6543](#). Software background; FORM 5 syntax is taken from the supplied manual.

- [14] The Sage Developers, *Finite Coxeter groups*, `degrees()` documentation. Definition as fundamental invariant degrees, including product groups and the repeated D_4 degree. [Official finite-Coxeter-group documentation](#).
- [15] A. D. Shapere and Y. Tachikawa, *Central charges of $\mathcal{N} = 2$ superconformal field theories in four dimensions*, JHEP 09 (2008) 109. [arXiv:0804.1957](#). Eqs. (2.5)–(2.6): free multiplets; §3 and Eq. (5.4): Casimir dimensions and $4(2a - c) = \sum_i (2\Delta_i - 1)$, with stated assumptions.
- [16] A. Gadde, L. Rastelli, S. S. Razamat and W. Yan, *Gauge Theories and Macdonald Polynomials*, Commun. Math. Phys. 319 (2013) 147–193. [arXiv:1110.3740](#). §4, Eqs. (4.5)–(4.11) and (4.15)–(4.17), Appendix B: HL and Coulomb limits.
- [17] Y. Tachikawa, *$\mathcal{N} = 2$ supersymmetric dynamics for pedestrians*, Lecture Notes in Physics 890 (2014). [arXiv:1312.2684](#). Lagrangian, hypermultiplet, flavor and moduli-space background.
- [18] C. Córdova, T. T. Dumitrescu and K. Intriligator, *Deformations of Superconformal Theories*, JHEP 11 (2016) 135. [arXiv:1602.01217](#). §3.2.2/Table 22: four-dimensional $\mathcal{N} = 2$ preserving deformations.
- [19] M. J. Kang, C. Lawrie, K.-H. Lee, M. Sacchi and J. Song, *Higgs, Coulomb, and Hall–Littlewood*, Phys. Rev. D 106 (2022) 106021. [arXiv:2207.05764](#). Examples and caveats showing that the HL index and Higgs Hilbert series need not coincide, including twisted class-S constructions.
- [20] The Stacks Project Authors, *Koszul regular sequences*, [Tag 062D](#). Regular sequences, Koszul homology and exactness criteria used in the constraint-counting derivation.
- [21] A. Hanany and R. Kalveks, *Highest Weight Generating Functions for Hilbert Series*, JHEP 10 (2014) 152. [arXiv:1408.4690](#). Representation-level organization of refined Hilbert series.